ThinPFN: a prior-fitted network for small-sample credit scoring, trained on one desktop GPU and validated on ten public portfolios

Javier Mancilla Montero
SquareOne Labs
javier.m@squareonecap.com

Abstract

Tabular foundation models replace the model-per-portfolio workflow of credit scoring with a single network, pre-trained on synthetic data, that reads a labelled table as context and scores new rows in one forward pass. We ask how much of that capability can be reproduced, understood and owned at small scale. ThinPFN is a 25.6M-parameter prior-fitted network pre-trained for six hours on one desktop GPU over 1.02 million synthetic tables drawn from structural causal priors and a credit ledger simulator; no real row enters pre-training. We evaluate it on ten public credit portfolios from four countries with a sealed test protocol: the test split of every book was read once, after the weights and the inference recipe were frozen, with contexts of 200, 500 and 1,000 labelled rows and paired bootstrap intervals against logistic regression, XGBoost tuned on the context, and TabPFN v2. ThinPFN ties TabPFN v2 on the ten-book mean (0.697 / 0.709 / 0.727 against 0.688 / 0.710 / 0.728 at the three context sizes), is ahead of it at the smallest context, beats tuned XGBoost in 23 of 30 cells (15 significantly, 2 significantly behind) and logistic regression in 29 of 30. Without any recalibration its probabilities have the lowest calibration error of the four models on validation splits and are never worse than the base-rate forecast on any book. Its weakness is concentrated in one regime, strong nonlinear signal, where boosted trees and the larger model stay three to six points ahead at small contexts. Two inference-time choices made with fixed weights, quantile-normal inputs and an average over column permutations, are worth 0.022 AUC, more than any change to the prior we tried. Code, weights, the benchmark with split manifests, and the training recipe are released for non-commercial use.

1 Introduction

In our experience of credit scoring work, the working answer to a tabular prediction problem in credit, fraud or collections has been a gradient-boosted tree ensemble such as XGBoost [2], fitted per portfolio, tuned per portfolio, and refitted every time a batch of outcomes arrives. The reasons are good ones: boosted trees are cheap, fast, handle missing values and mixed types, and the lenders we have worked with have found their explanation tooling acceptable to their model-risk functions. Benchmarks that compare them with neural networks on medium-sized tables keep finding them ahead or tied [3, 11]. The cost of this regime is not accuracy; it is the unit of work. Every new product, segment, country or partner needs its own model, its own validation cycle and its own retraining schedule, and the first months of a new portfolio, when a few hundred outcomes exist, are exactly when a dedicated model is least reliable.

Tabular foundation models change the unit of work. A network is pre-trained once, on synthetic tables, to perform Bayesian prediction in context: the labelled rows of a new problem are the prompt, the unlabelled rows are the query, and the prediction is the posterior predictive of a prior the network has absorbed during pre-training [5, 12]. No gradient step happens at decision time. In principle a new portfolio can be scored as soon as its first few hundred outcomes are known, and a confirmed outcome can enter the next decision as soon as it is appended to the context; whether a lender's controls allow that cadence is an operational question this paper does not study. The second generation of these models [6, 15] reaches or beats tuned boosting on standard benchmarks up to tens of thousands of rows.

The published models were pre-trained with budgets well beyond the six GPU-hours used here (their papers report the details [6, 15]) and are released as weights. We ask a narrower and, for a practitioner, more use-

*ThinPFN is a personal research project of the author, published from the author's GitHub account at https://github.com/maximer-v/thinpfn. The author is affiliated with SquareOne Labs.

ful question: how much of this capability can one person reproduce, understand and own, with one desktop GPU and a few hours of training, and does a prior built around the mechanics of credit buy anything on real portfolios? The question matters because ownership is the point. A team that can regenerate its prior, retrain its model and audit every line between the synthetic generator and the score has a different relationship with the model than one that downloads weights.

This paper makes three contributions. First, a complete, small instance of a tabular foundation model, ThinPFN, with its synthetic prior, network, pretraining recipe and inference recipe documented and released. Second, a benchmark of ten public credit portfolios from four countries with shipped split manifests, baselines fitted on identical contexts, paired bootstrap intervals, and a sealed test split read once. Third, a measured account of what moved the result and what did not, including a negative one: making the synthetic stream statistically resemble real books did not help, while depth, compute and two inference-time choices did.

2 Background

Prior-fitted networks. A prior-fitted network (PFN) is trained on tasks sampled from a prior over data-generating processes. Each task is a table with labelled context rows and unlabelled query rows; the network is trained to predict the query labels given the context, by cross-entropy over millions of such tasks. Because the training distribution is the prior, the network approximates the posterior predictive under that prior for any new table, in one forward pass [12]. TabPFN [5] applied this to small tabular classification with a prior of random structural causal models (SCMs) and Bayesian neural networks; the second version [6] added a per-cell architecture that attends across both rows and features, a richer prior, and support for regression and larger tables. TabICL [15] separates column embedding, row interaction and in-context learning into stages so that the in-context stage sees one token per row, which scales the approach to hundreds of thousands of rows. TabDPT [10] pre-trains the same kind of network on real tables instead of a synthetic prior.

The ecosystem around them. We find it useful to group the "foundation" approaches on offer for tabular data into three informal families; the grouping is ours, not a published taxonomy. Embedding approaches train a transformer to produce a vector per row and feed that vector to a conventional classifier; they keep the second model and its retraining. Relational approaches operate over several linked tables at once and are strongest where the schema itself carries signal. Single-table in-context models, the family of TabPFN and TabICL, score directly from the labelled table and need no second model; they are the natural fit for a scoring team whose data is already a feature table with a label. ThinPFN belongs to the third family. The operational case for it is a hypothesis rather than a measured result: one model per family of problems instead of per portfolio, and decisions that can be refreshed as outcomes arrive. Model-agnostic explanation methods (perturbation-based reason codes, Shapley attribution, distillation into a simpler surrogate) apply to any scorer and so to this one; we have not evaluated them on ThinPFN and make no claim about their adequacy for a given regulator.

Why synthetic priors. The prior is the only place where domain knowledge enters a PFN. A generic SCM prior encodes the belief that features are produced by sparse causal mechanisms with noise; it is agnostic about what the columns mean. A domain prior can encode mechanics: that defaults arise from a delinquency process that evolves month by month, that a bureau snapshot summarises that process with lag, that an application form measures capacity and stability with error, and that the label is a function of the process on a horizon. Whether such mechanics transfer to real portfolios, where the columns are whatever the lender happened to record, is an empirical question. We answer it in Section 5.

3 ThinPFN

3.1 The prior

Every training table is generated, never collected. A table is drawn from one of three generators.

Structural causal priors. A causal graph over k latent causes and d observed columns is sampled; the mechanism is either a random multilayer perceptron with sampled depth, width, activation and noise, or a random ensemble of decision trees with sampled depth and count. A continuous target is produced by the same mechanism and thresholded at a sampled prevalence to give a binary label. Hyper-parameters are drawn from log-uniform ranges so that most tables are moderately learnable and a few are hard.

Credit ledger simulator. A portfolio is drawn: a product tag (card, instalment, peer-to-peer, vehicle; the tag is metadata and does not change the mechanics), a prevalence between 2 and 35 percent, an observation window of three to eleven months, a default horizon of one to twelve months, a label definition (ever 30, 60 or 90 days late, or the end state), a thin-file share, a

drift strength and noise levels. Each borrower has latent traits whose linear combination is a risk score. A Markov delinquency chain over eleven states (no consumption, paid in full, revolving, one to eight months late) evolves monthly; the probability of slipping one state rises with the risk score and with a smooth random macro path, the probability of curing falls with it, and the 90-plus states are nearly absorbing. The label is one if the chain meets the label definition within the horizon after the observation window; prevalence is matched by bisection on the slip intercept. Observed columns are emitted in blocks, each present with its own probability: the ledger (monthly delinquency status in the coding of the UCI Taiwan dataset, and optionally bill and payment amounts), a credit limit and application columns (noisy monotone transforms of single traits), bureau columns (noisy monotone transforms of the risk score), utilisation and payment ratios derived from the ledger, nominal demographic columns and an age, and an application date when the portfolio drifts. A share of borrowers are thin files and lose the ledger block.

Drift. With a small probability the generic tables receive a cohort index in $[0, 1]$ and are generated in cohort buckets: for the MLP prior a sparse subset of the mechanism's weights moves along the cohort; for the tree prior each bucket draws its own tree ensemble, so the mechanism changes between cohorts rather than shifting smoothly; with probability one half the continuous target is also tilted along the cohort (prevalence drift). The context can then be the early cohorts and the query the late ones.

Every table passes the same contract: a random context share between 0.3 and 0.9, missing values injected completely at random on a few columns or dependent on the value of one column, nuisance columns (noise, collinear copies, ratios), standardisation of numeric columns on the context rows, and a learnability filter that drops tables on which an extra-trees classifier fitted to the context scores below 0.55 or above 0.995 AUC on the query. The released mixture is 0.40 MLP-SCM, 0.20 tree-SCM, 0.30 ledger, 0.07 drifting MLP-SCM and 0.03 drifting tree-SCM, with 256 to 1,024 rows and up to 64 columns per table.

3.2 The network

ThinPFN is a per-row model. Each row becomes one token: a linear map of its standardised feature vector concatenated with a per-cell missingness indicator and a per-column categorical flag (the released weights were trained with the flag off), plus a learned label embedding on context rows and an "unknown" embedding on query rows, plus a periodic embedding of the cohort index. Twenty-four learned "thinking" rows are prepended. A twelve-layer pre-norm transformer of width 512 with eight heads runs over the rows under a mask in which context rows and thinking rows attend to each other and to themselves, and query rows attend to the context, the thinking rows and themselves. A query row's score therefore never depends on which other query rows are in the batch, and scoring can be batched freely. A linear head produces two logits per row; the loss is the cross-entropy on query rows. The model has 25.6M parameters. The column-wise and row-wise attention stages of the staged architectures [15] are implemented but were too slow to train at this budget; a 3.3M per-row model with the same prior reaches the same development accuracy as the staged model at a fraction of the cost, so the released model has no such stages.

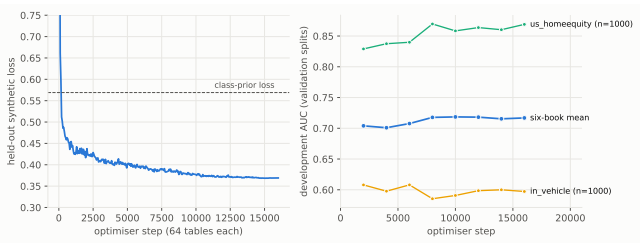


Figure 1: Pre-training. Left: loss on a fixed held-out synthetic batch against the class-prior loss it must beat. Right: development AUC (validation splits, never the test split) on the six charter books, with the two books that moved most.

3.3 Pre-training

The released weights are one run of 16,000 AdamW steps with 64 tables per step (1.02 million tables), learning rate 5×10^{-4} with 400 warm-up steps and cosine decay to zero, weight decay 0.01, gradient clipping at 1.0, bf16 autocast, and tables generated on the fly by twelve CPU workers. The run took six hours on a single desktop GPU with unified memory; it is bound by table generation, not by the network, which holds five gigabytes. A fixed held-out synthetic batch is scored every fifty steps; the last six snapshots, 500 steps apart, are averaged into the released weights. The real-book development AUC tracked the held-out synthetic loss throughout (Figure 1), which is the practical signal that the prior is being fitted rather than memorised.

3.4 Inference

Two choices made after pre-training, on validation splits and with the weights fixed, define the released inference recipe. First, numeric columns are transformed to a quantile-normal scale fitted on the context rows, instead of the z-score clipped at four standard deviations that the prior was standardised with. Real credit columns are heavy-tailed and discrete (half the columns of the

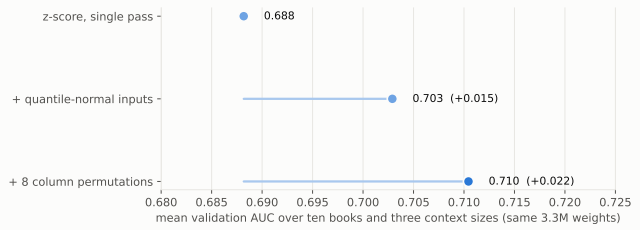


Figure 2: Inference recipe on fixed weights (3.3M development model, validation splits, ten books, three context sizes).

benchmark books have fewer than six percent distinct values; money columns carry point masses); the z-score squashes most of their rows into a narrow band while the prior's own inputs are close to Gaussian. Second, the log-odds are averaged over eight random column orders. The row embedding is a slot-wise linear map, so the network is not permutation-invariant over columns, and the orders act as a cheap ensemble. On the 3.3M development model the two changes together are worth 0.022 AUC on the ten-book mean, with no cell significantly worse (Figure 2); the released model inherits both.

4 Benchmark

Books. Ten public consumer-credit portfolios with a binary default label (Table 1). Six form the charter set the study was designed around; four were added for breadth from lenders in the United States, China and India. Two pairs share an origin (the two Taiwan views; the two LendingClub extracts, from different periods and with different fields; the two Taiwan views share their 4,500 test rows), so the set contains eight source datasets. Each book carries a frozen four-way split into training, validation, calibration and test rows; two books with origination dates (LendingClub, Prosper) are split out of time, the rest at random with stratification. The two Taiwan books are behavioural scoring (six months of repayment history, label the next month); the other books use origination-time fields, with fields observed after origination (payments received, recoveries, current balance) excluded, and PPDai's 29 fields taken as released. String columns are ordinal-coded on the whole source table (label-free), a transductive step the sealed test inherits. TVS contains many identical feature rows across splits (1,919 of the test rows match a training row exactly); the bootstrap unit is the row everywhere. The repository ships the split manifests keyed to the rows of the public source files, whose hashes are recorded, and a preparation script that rebuilds every table.

Data sources. The eight source datasets, with the record each book was built from (all accessed 2026-11-09; the repository records the exact URL and SHA-256 of every raw file). Taiwan credit cards: the UCI record [17] of the data introduced by Yeh and Lien [18]. Give Me Some Credit: the 2011 Kaggle competition [7], read from its OpenML copy. L&T vehicle loans: the 2019 LTFS FinHack data [9], read from OpenML. HMEQ: the home equity dataset that accompanies Baesens et al. [1], read from OpenML. LendingClub: the company's public loan files [8], in two extracts (an OpenML 2007 to 2015 subset for `us_p2p`; the 2007 to 2018Q4 file, restricted to 2007 to 2015 vintages, for `lending_club`). PPDai: the figshare release [4]. Prosper: the Udacity copy of the listings file [14]. TVS Credit: the OpenML upload [13]. OpenML copies are served by the platform of Vanschoren et al. [16]. Field selections, label definitions, split rules and the post-origination exclusions are ours and are stated per book in the repository; period and licence statements are those of the records cited.

Protocol. For each book and each context size $n \in \{200, 500, 1000\}$, three contexts are drawn from the training split with stratified sampling (seeds 0, 1, 2). Every model scores the same query rows (the evaluation split, subsampled to 4,000 rows when larger) from the same context. The metric is ROC AUC, averaged over the three seeds. The difference between ThinPFN and each baseline carries a 95% paired bootstrap interval over the query rows (1,000 resamples, seeds averaged inside each resample); a cell is "significant" when the interval excludes zero.

Baselines. Logistic regression with median imputation and standardisation. XGBoost with a small grid of depth, learning rate and regularisation chosen on a held-out slice of the context with early stopping, which is how a practitioner would tune it with this little data; at $n < 300$ a fixed shallow configuration is used. TabPFN v2 [6] with its released weights and default settings.

Sealed test. Development used validation splits only, many times. The test split of every book was read once, by a dedicated script, after the released weights and the inference recipe were frozen. Three claims were written before that read: G2, the original gate of the study, that ThinPFN beats both TabPFN v2 and XGBoost on the six charter books with the interval above zero on at least four and never significantly behind; C1, that at $n = 200$ it is significantly ahead of both on at least four of six; and C2, that the ten-book mean is within 0.01 of TabPFN v2 at $n = 500$ and $n = 1000$. The results below report all three.

Table 1: The ten books. Rows are those used after the split rules; prevalence is of the test split; OOT = out-of-time splits by origination date.

book	origin	rows	feat.	prev.
uci_full	UCI 350, Taiwan cards, 2005	30,000	23	0.22
tw_cards	same table, 4,000 training rows	16,000	23	0.22
us_bureau	Give Me Some Credit, US	9,335	10	0.06
in_vehicle	L&T vehicle loans, India	9,335	19	0.22
us_homeequity	HMEQ home equity, US	5,960	12	0.20
us_p2p	LendingClub 2007 to 2015	7,832	13	0.16
lending_club	LendingClub 2007 to 2015 vintages, OOT	642,138	28	0.15
ppdai	PPDai, China	55,596	29	0.13
prosper	Prosper 2005 to 2014, OOT	38,907	36	0.15
tvs	TVS Credit cross-sell, India	119,528	29	0.02

Table 2: Sealed test: mean ROC AUC over ten books.

n	ThinPFN	TabPFN v2	XGBoost	Logistic
200	**0.697**	0.688	0.654	0.654
500	0.709	0.710	0.674	0.665
1000	0.727	0.728	0.706	0.679

Cells ahead of 30 (sig. ahead / behind): TabPFN v2 17 (7/5), XGBoost 23 (15/2), logistic 29 (23/0)

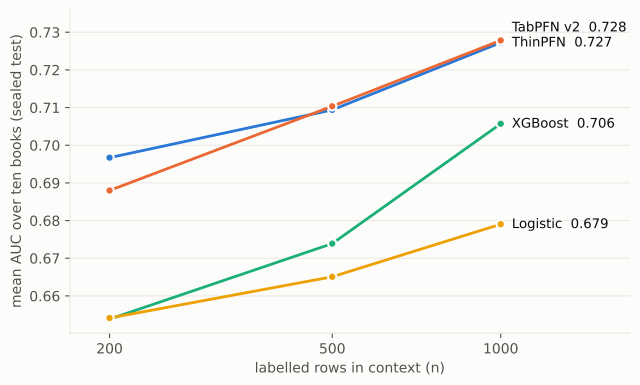


Figure 3: Sealed test: mean AUC over the ten books against context size.

5 Results

Main result. Table 2 and Figure 3 give the ten-book means. ThinPFN ties TabPFN v2 on the mean, is ahead of it by 0.009 at $n = 200$ and within 0.001 at the two larger contexts, and is ahead of tuned XGBoost by 0.043, 0.035 and 0.021. Against logistic regression it is ahead in 29 of 30 cells (the exception is in_vehicle at $n = 200$, -0.010). C2 passes. G2 and C1 fail: at $n = 200$ the interval is above zero against both baselines on two of the six charter books (tw_cards and us_p2p), not four, and one charter book is significantly behind both. ThinPFN does not have a decisive win over TabPFN v2; it has parity at a training budget of six desktop GPU-hours, a small-context advantage on several books, and a documented failure mode.

Per book. Figure 4 shows every cell. The significant wins over TabPFN v2 at $n = 200$ are LendingClub ($+0.062$), tw_cards ($+0.038$) and us_p2p ($+0.016$); in_vehicle is ahead at every n but significant only at $n = 500$. The losses are concentrated in one book, us_homeequity, at -0.039, -0.038 and -0.028, which is also the only book where XGBoost is significantly ahead of ThinPFN. The remaining significant losses to TabPFN v2 are uci_full at $n = 500$ (-0.009) and tw_cards at $n = 1000$ (-0.013). The XGBoost comparison carries a caveat: on TVS (2 percent prevalence, four positives in a 200-row context) the small-context XGBoost configuration collapses to a constant score for every seed at $n = 200$ and one seed at $n = 500$, so those cells (AUC 0.50) inflate the count; excluding TVS, ThinPFN is ahead of XGBoost in 20 of 27 cells (12 significantly, 2 significantly behind), with mean differences of $+0.019$, $+0.009$ and $+0.008$. On the three weak-signal books, where no model reaches 0.66 AUC (in_vehicle, ppdai, us_p2p), ThinPFN is ahead of TabPFN v2 in seven of nine cells and of XGBoost in eight of nine.

What explains the losses. On the development splits, ThinPFN's deficit to TabPFN v2 at $n = 1000$ correlates -0.87 with the book's achievable AUC (a tuned boosting model on the full training split), -0.73 with the AUC of the single best feature, and -0.69 with the nonlinearity gap (boosting minus logistic). The model regularises well where there is little to find and underexploits strong nonlinear structure. us_homeequity has a nonlinearity gap of 0.13, the largest in the set; a twelve-layer per-row network reaches 0.89 there where the larger model reaches 0.92. The development evidence below, where the gains from a deeper and longer-trained model landed on exactly these books, points to network capacity rather than to the prior; since that comparison

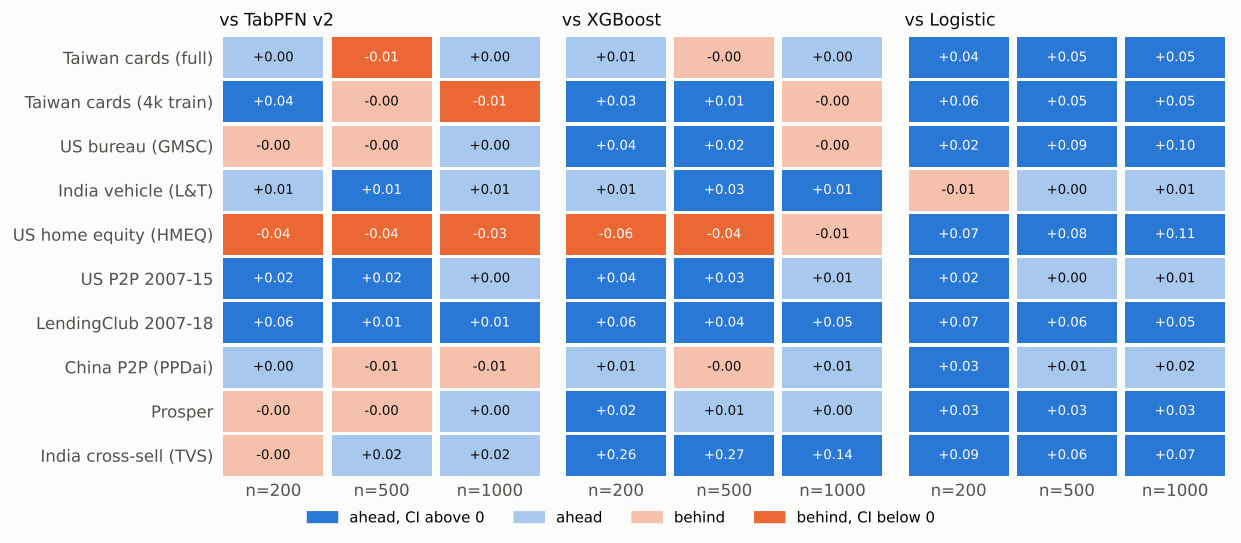


Figure 4: Every cell of the sealed test. Colour: sign of the AUC difference of ThinPFN against the baseline, saturated when the 95% paired bootstrap interval excludes zero.

changed depth and training budget together, capacity is the working hypothesis, not an established cause.

5.1 Probability quality

A credit score is consumed as a probability of default, so ranking is half of the question. After the sealed test we added a probability study with the same contexts and baselines, scored on the validation split of every book: the test split is not read a second time, and since the validation splits were used during development to choose the inference recipe by AUC, the study is descriptive and carries no pre-written claim. No model is recalibrated. We report the Kolmogorov-Smirnov statistic (KS), the Brier skill against the constant forecast of the context's default rate, and the expected calibration error over ten equal-count bins (ECE), each averaged over three seeds and then over the ten books (Table 3, Figure 5).

ThinPFN has the lowest mean calibration error at every context size, lower than TabPFN v2 on nine of ten books at each n, and it is the only model with a positive Brier skill on every book at every n, including $n = 200$, where TabPFN v2 falls below the base rate on two books, XGBoost on four and logistic regression on six. TabPFN v2 separates slightly better: it leads on mean KS at every n and on mean Brier skill at $n = 500$ and $n = 1000$. Most of that Brier gap is us_homeequity (0.338 against 0.451 at $n = 500$), the book that carries the AUC losses; book by book ThinPFN has the higher skill in 20 of 30 cells. The predicted level is anchored on the context: the mean prediction of ThinPFN exceeds the observed default rate by about half a point on average (0.175 against 0.169), and on the out-of-time books part of the error of every model is the movement of the default rate between the context period and the scored period (LendingClub: 0.170 predicted, 0.147 observed), which no model can read from the context. These are means without intervals; the differences in calibration error are small and were not tested for significance.

Table 3: Probability quality on validation splits, mean of ten books, no recalibration. Skill: Brier skill against the base-rate forecast. Last column: books where the model is worse than the base rate (negative skill).

n	model	KS	skill	ECE	worse
200	ThinPFN	0.315	**0.088**	**0.025**	0
	TabPFN v2	**0.321**	0.073	0.032	2
	XGBoost	0.259	0.074	0.037	4
	Logistic	0.257	−0.023	0.068	6
500	ThinPFN	0.339	0.104	**0.023**	0
	TabPFN v2	**0.351**	**0.113**	0.029	0
	XGBoost	0.297	0.095	0.028	0
	Logistic	0.285	0.032	0.043	4
1000	ThinPFN	0.365	0.119	**0.023**	0
	TabPFN v2	**0.373**	**0.128**	0.027	0
	XGBoost	0.340	0.120	0.024	0
	Logistic	0.307	0.071	0.032	1

Latency. One scoring call with the default eight column orders, 1,000 context rows and 2,000 query rows of 20 features, takes 5.7 s on an Arm CPU limited to eight threads and 0.28 s on the desktop GPU used for pre-training; a single pass takes 0.75 s and 0.04 s. A single query against 1,000 context rows takes 1.0 s and 0.09 s with the ensemble, since every call reads the con-

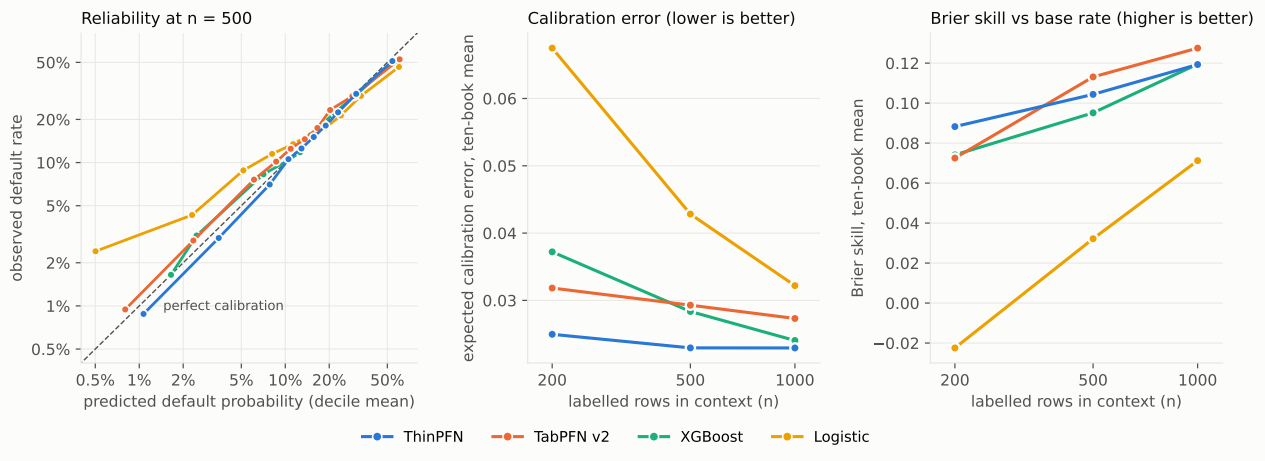


Figure 5: Probability quality without recalibration (validation splits, ten books). Left: reliability at $n = 500$, rows pooled over books and seeds, deciles of the predicted probability. Centre and right: expected calibration error and Brier skill against the base-rate forecast, mean of ten books.

text again; a context and a query of 4,000 rows each take 29 s and 1.2 s.

5.2 What moved the result during development

Four findings from the validation splits, reported here because they are the transferable part of the work. The prior ablations were run at equal compute (3.3M model, 160k tables); the size comparison was not. The development record behind them (realism audit, leave-dataset-out runs, prior ablations) is summarised rather than released; the released result files are the sealed test, the probability study, the inference study and the size comparison.

Depth and compute. With the same prior and the same inference recipe, a 3.3M model trained on 160k tables reaches 0.700 / 0.712 / 0.724 on the ten-book validation mean; the released 25.6M model on 1.02M tables reaches 0.706 / 0.721 / 0.737 (Figure 6). The gain is concentrated on clean, strongly nonlinear books (us_homeequity at $n = 1000$: 0.848 to 0.904) and costs under a point on the weakest book. Within every run, the real-book AUC followed the synthetic held-out loss, so the lever is fitting the prior better. The released run was under-driven (gradient norms of 0.1 to 0.2 against a clip of 1.0, at half the learning rate of the small model); a higher learning rate is the first untested training lever.

Statistical realism of the prior did not help. An audit of 25 per-table statistics on 330 real tables against the synthetic stream showed the stream far from real books on every axis (labels too predictable, columns too continuous, too few categoricals, too symmetric). A second simulator with product and feature families and realistic marginals, plus discretisation and prevalence-preserving

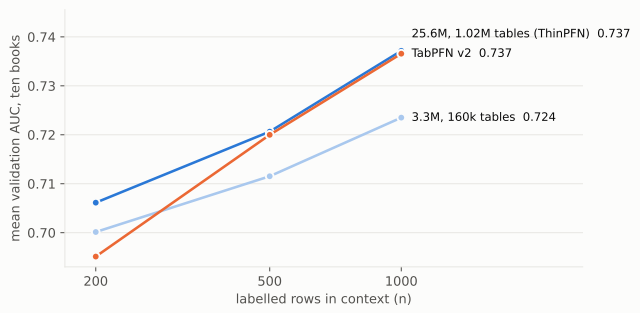


Figure 6: Same prior, same inference recipe, two model sizes (validation splits).

label noise applied to every component, brought the stream within noise of real books on most statistics (a leave-dataset-out discriminator fell from 0.997 to 0.909). Trained at equal compute, the model was worse on the mean at every n and significantly worse on four to six of ten books. Label noise in training appears to teach the network to trust the context less, which hurts at larger n. The prior that works is the simple one, and the lever against mismatch is on the inference side (Section 3.4).

Real rows in the training stream hurt. Mixing subtables of the public books (never the held-out lender) into the stream, or copula twins of them, or a field-type vocabulary through which knowledge could transfer across lenders, failed to beat the synthetic-only model on an unseen lender in nine of nine leave-dataset-out runs. At this scale, real data belongs in evaluation, not in pre-training.

The ledger component is a trade. At equal compute, a generic-only prior beat the released mixture on the Taiwan books (uci_full at $n = 1000$: +0.015) and lost on the weakest book (in_vehicle: −0.02). The credit

mechanics buy the weak-signal edge and cost clean-book accuracy; the weight is a choice, not a free parameter.

6 Limitations and intended use

ThinPFN is for the small-sample regime: contexts of 200 to a few thousand rows and at most 64 numeric or ordinal features, binary labels. It was pre-trained on tables of at most 1,024 rows and stops improving above a few thousand context rows, where a tuned boosted model with all the data is the better tool. On books with strong nonlinear signal it trails TabPFN v2 by three to four points and XGBoost by four to six points at $n \leq 500$. The inference recipe and the depth were chosen on validation splits that were read many times; the sealed test was read once, and the significant wins over TabPFN v2 at $n = 200$ shrank from five books on validation to three on test, which is the expected size of that selection effect. The released weights are one pre-training seed; variance across seeds has not been measured. Predicted probabilities are anchored on the default rate of the context rows: a context sampled at a different rate than the scored population, or a shift of the default rate over time, moves the calibration and needs a correction outside the model, and the probability study of Section 5.1 rests on validation splits, not on the sealed test. The model is released for research and non-commercial use; a lender considering any automated decision about people must satisfy, independently of the model, whatever its own regulation requires of fairness assessment, human review and explanations.

7 Reproducibility

The repository contains the package (prior, network, trainer, scikit-learn style front end), the released weights (fp16, 51 MB), the benchmark with split manifests and a preparation script that rebuilds every book from its public source, the sealed-test script and its output, the probability study and latency scripts with their outputs, the figure script, tests, and a walkthrough notebook. The pre-training command that produced the released weights is one line; the benchmark is two. The sealed-test numbers, the probability study, the latency table, the inference study, the size comparison and the training log are in `benchmark/results/`; the other development findings are reported from the development record. A resumed training run restarts the synthetic streams, so only an uninterrupted run reproduces the released weights exactly. The datasets keep the terms stated at the records cited in the data-sources paragraph; one (Prosper) has no stated licence and is used for evaluation only.

References

[1] Bart Baesens, Daniel Roesch, and Harald Scheule. *Credit Risk Analytics: Measurement Techniques, Applications, and Examples in SAS*. Wiley, 2016. The HMEQ home equity dataset accompanies the book at `http://www.creditriskanalytics.net`; copy used: OpenML dataset 43337 (CC0), accessed 2026-11-09.

[2] Tianqi Chen and Carlos Guestrin. XGBoost: A scalable tree boosting system. In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 785–794, 2016. doi: 10.1145/2939672.2939785.

[3] Léo Grinsztajn, Edouard Oyallon, and Gaël Varoquaux. Why do tree-based models still outperform deep learning on typical tabular data? In *Advances in Neural Information Processing Systems 35, Datasets and Benchmarks Track*, 2022. URL `https://papers.nips.cc/paper_files/paper/2022/hash/0378c7692da36807bdec87ab043cdadc-Abstract-Datasets_and_Benchmarks.html`.

[4] Hongliang He. PPDaiData. figshare dataset, CC BY 4.0, `https://doi.org/10.6084/m9.figshare.13573856`, 2021.

[5] Noah Hollmann, Samuel Müller, Katharina Eggensperger, and Frank Hutter. TabPFN: A transformer that solves small tabular classification problems in a second. In *International Conference on Learning Representations*, 2023.

[6] Noah Hollmann, Samuel Müller, Lennart Purucker, Arjun Krishnakumar, Max Körfer, Shi Bin Hoo, Robin Tibor Schirrmeister, and Frank Hutter. Accurate predictions on small data with a tabular foundation model. *Nature*, 637:319–326, 2025. doi: 10.1038/s41586-024-08328-6.

[7] Kaggle. Give Me Some Credit. Prediction competition, `https://www.kaggle.com/c/GiveMeSomeCredit`; copy used: OpenML dataset 46929 (TabArena curation), accessed 2026-11-09, 2011.

[8] LendingClub Corporation. Loan data, 2007 to 2018Q4. Public loan statistics files, discontinued by the company in 2020; copies used: OpenML dataset 43729 (CC0, 2007 to 2015 subset) and the Hugging Face mirror `https://huggingface.co/datasets/codesignal/lending-club-loan-accepted`, accessed 2026-11-09, 2019.

[9] L&T Financial Services and Analytics Vidhya. LTFS data science FinHack: vehicle loan default prediction. Hackathon dataset, `https://data`

hack.analyticsvidhya.com/contest/ltfs-datascience-finhack-an-online-hackathon/; copy used: OpenML dataset 46430, accessed 2026-11-09, 2019.

[10] Junwei Ma, Valentin Thomas, Rasa Hosseinzadeh, Alex Labach, Jesse C. Cresswell, Keyvan Golestan, Guangwei Yu, Anthony L. Caterini, and Maksims Volkovs. TabDPT: Scaling tabular foundation models on real data. In *Advances in Neural Information Processing Systems 38*, 2025. URL https://papers.nips.cc/paper_files/paper/2025/hash/fc0e3f908a2116ba529ad0a1530a3675-Abstract-Conference.html.

[11] Duncan McElfresh, Sujay Khandagale, Jonathan Valverde, Vishak Prasad C., Ganesh Ramakrishnan, Micah Goldblum, and Colin White. When do neural nets outperform boosted trees on tabular data? In *Advances in Neural Information Processing Systems 36, Datasets and Benchmarks Track*, 2023. URL https://arxiv.org/abs/2305.02997.

[12] Samuel Müller, Noah Hollmann, Sebastian Pineda Arango, Josif Grabocka, and Frank Hutter. Transformers can do Bayesian inference. In *International Conference on Learning Representations*, 2022.

[13] OpenML contributors. TVS_Loan_Default. OpenML dataset 43743 (CC0), uploaded 2022-03-24, TVS Credit cross-sell personal loans, India; accessed 2026-11-09, 2022.

[14] Prosper Marketplace. Prosper loan data, 2005 to 2014. Copy distributed by Udacity for course ud651, https://s3.amazonaws.com/udacity-hosted-downloads/ud651/prosperLoanData.csv, no stated licence, accessed 2026-11-09, 2014.

[15] Jingang Qu, David Holzmüller, Gaël Varoquaux, and Marine Le Morvan. TabICL: A tabular foundation model for in-context learning on large data. In *Proceedings of the 42nd International Conference on Machine Learning, PMLR 267*, 2025.

[16] Joaquin Vanschoren, Jan N. van Rijn, Bernd Bischl, and Luís Torgo. OpenML: networked science in machine learning. *SIGKDD Explorations*, 15(2):49–60, 2014. doi: 10.1145/2641190.2641198.

[17] I-Cheng Yeh. Default of credit card clients. UCI Machine Learning Repository, dataset 350, https://doi.org/10.24432/C55S3H, 2009.

[18] I-Cheng Yeh and Che-hui Lien. The comparisons of data mining techniques for the predictive accuracy of probability of default of credit card clients. *Expert Systems with Applications*, 36(2):2473–2480, 2009. doi: 10.1016/j.eswa.2007.12.020.